

# Assignment 2

## Graduate Admission

**This dataset is created for prediction of Graduate Admissions from an Indian perspective.**

The dataset contains several parameters which are considered important during the application for Masters Programs. The parameters included are :

- GRE Scores ( out of 340 )
- TOEFL Scores ( out of 120 )
- University Rating ( out of 5 )
- Statement of Purpose and Letter of Recommendation Strength ( out of 5 )
- Undergraduate GPA ( out of 10 )
- Research Experience ( either 0 or 1 )
- Chance of Admit ( ranging from 0 to 1 )

## Import the Required Libraries

In [1]:

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import random
```

In [2]:

```
df = np.genfromtxt('Admission_Predict.csv', delimiter = ',', dtype=str);
df
```

Out[2]:

```
array([[ 'Serial No.', 'GRE Score', 'TOEFL Score', ..., 'CGPA',
        'Research', 'Chance of Admit'],
       ['1', '337', '118', ..., '9.65', '1', '0.92'],
       ['2', '324', '107', ..., '8.87', '1', '0.76'],
       ...,
       ['398', '330', '116', ..., '9.45', '1', '0.91'],
       ['399', '312', '103', ..., '8.78', '0', '0.67'],
       ['400', '333', '117', ..., '9.66', '1', '0.95']], dtype='<U17')
```

In [3]:

```
headers = df[0,1:]; # TO not take serial no
print(headers)
data = np.array(df[1:,1:], dtype=float); # This will take from the GRE Score
print(data)
```

```
['GRE Score' 'TOEFL Score' 'University Rating' 'SOP' 'LOR ' 'CGPA'
 'Research' 'Chance of Admit']
[[337.  118.    4.   ...   9.65   1.    0.92]
 [324.  107.    4.   ...   8.87   1.    0.76]
 [316.  104.    3.   ...    8.     1.    0.72]
 ...
 [330.  116.    4.   ...   9.45   1.    0.91]
 [312.  103.    3.   ...   8.78   0.    0.67]
 [333.  117.    4.   ...   9.66   1.    0.95]]
```

In [4]:

```
data_norm = (data-np.mean(data, axis = 0))/np.std(data, axis = 0)
```

In [5]:

```
# Extract y from data

y_label = 'Chance of Admit';
y_index = np.where(headers == y_label)[0][0];
y = data_norm[:,y_index];

# Extract x from data

X = data_norm[:,0:y_index];
```

In [6]:

```
# Insert column of 1's for intercept column
X = np.insert(X, 0, 1, axis=1) # added the intercept
```

In [7]:

```
print(X.shape)
```

```
(400, 8)
```

In [8]:

```
print(y.shape)
```

```
(400,)
```

In [9]:

```
print(X[0])
```

```
[1.          1.76210664 1.74697064 0.79882862 1.09386422 1.16732114
 1.76481828 0.90911166]
```

In [10]:

```
m = X.shape[0]
n = X.shape[1]
```

In [11]:

```
# Partion data into training and test datasets

idx = np.arange(0,m)
random.shuffle(idx)

percent_train = .6
m_train = int(m * percent_train)
train_idx = idx[0:m_train]
test_idx = idx[m_train:m+1]
X_train = data[train_idx,1:y_index];
X_test = data[test_idx,1:y_index];

y_train = data[train_idx,y_index];
y_test = data[test_idx,y_index];
```

In [12]:

```
# Cost function normalized by number of examples

def H(theta,X,y):
    return 1 / 2 / X.shape[1] * (h(X,theta)-y).T.dot(h(X,theta)-y)
```

In [13]:

```
# Solve the normal equations

def regress(X, y):
    cov = np.dot(X.T, X)
    cov_inv = np.linalg.inv(cov)
    theta = np.dot(cov_inv, np.dot(X.T, y))
    return theta
regress(X, y)
```

Out[13]:

```
array([[ 8.70408612e-16,  1.39783600e-01,  1.24258432e-01,  4.58476504e-02,
        -2.33355774e-02,  1.40830779e-01,  4.97342141e-01,  8.57053356e-02])
```

In [14]:

```
# Cost function normalized by number of examples

def J(theta, X, y):
    return 1 / 2 / X.shape[1] * (h(X, theta) - y).T.dot(h(X, theta) - y)

# Get design matrix for polynomial model of degree d

def x_polynomial(x, d):
    a = np.ones((x.shape[0], 1))
    for i in range(d):
        a = np.concatenate((a, x**(i+1)), axis = 1)
    return a
```

In [15]:

```
# Build models of degree 1 to max_degree

max_degree = 2

J_train = np.zeros(max_degree)
J_test = np.zeros(max_degree)
```

## Split the data into training and tests

In [16]:

```
# splits the training and test data set in 80% : 20%
# assign random_state to any value. This ensures consistency.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=
5)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(320, 8)
(80, 8)
(320,)
(80,)
```

## Linear Regression

In [17]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

lin_model = LinearRegression()
lin_model.fit(X_train, y_train)
```

Out[17]:

LinearRegression()

In [18]:

```
# model evaluation for training set

y_train_predict = lin_model.predict(X_train)
rmse = (np.sqrt(mean_squared_error(y_train, y_train_predict)))
r2 = r2_score(y_train, y_train_predict)

print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")

# model evaluation for testing set

y_test_predict = lin_model.predict(X_test)
# root mean square error of the model
rmse = (np.sqrt(mean_squared_error(y_test, y_test_predict)))

# r-squared score of the model
r2 = r2_score(y_test, y_test_predict)

print("The model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
```

The model performance for training set

-----

RMSE is 0.4535600939850745

R2 score is 0.8069022125298158

The model performance for testing set


-----

RMSE is 0.40541316382029674

R2 score is 0.772829262603884

In [19]:

```
# plotting the y_test vs y_pred
# ideally should have been a straight line
plt.scatter(y_test, y_test_predict)
plt.show()
```



## Polynomial Regression

Let's apply the Polynomial Regression with degree 2 and test.

To generate the higher order degrees, we use PolynomialFeatures class from sklearn library.

In [22]:

```
from sklearn.preprocessing import PolynomialFeatures

def polynomial_regression_model(degree):
    "Creates a polynomial regression model for the given degree"
    poly_features = PolynomialFeatures(degree=degree)

    # transform the features to higher degree features.
    X_train_poly = poly_features.fit_transform(X_train)

    # fit the transformed features to Linear Regression
    poly_model = LinearRegression()
    poly_model.fit(X_train_poly, y_train)

    # predicting on training data-set
    y_train_predicted = poly_model.predict(X_train_poly)

    # predicting on test data-set
    y_test_predict = poly_model.predict(poly_features.fit_transform(X_test))

    # evaluating the model on training dataset
    rmse_train = np.sqrt(mean_squared_error(y_train, y_train_predicted))
    r2_train = r2_score(y_train, y_train_predicted)

    # evaluating the model on test dataset
    rmse_test = np.sqrt(mean_squared_error(y_test, y_test_predict))
    r2_test = r2_score(y_test, y_test_predict)

    print("The model performance for the training set")
    print("-----")
    print("RMSE of training set is {}".format(rmse_train))
    print("R2 score of training set is {}".format(r2_train))

    print("\n")
    print("The model performance for the test set")
    print("-----")
    print("RMSE of test set is {}".format(rmse_test))
    print("R2 score of test set is {}".format(r2_test))
```

In [23]:

```
polynomial_regression_model(2)
```

```
The model performance for the training set
-----
RMSE of training set is 0.42904975210889146
R2 score of training set is 0.8272082770286167
```

```
The model performance for the test set
-----
RMSE of test set is 0.4071697317313779
R2 score of test set is 0.7708564341794006
```

In [ ]: